


# Google Cloud Functions


# Cloud Functions



- **Imagine you want to execute some code when an event happens?**
  - A file is uploaded in Cloud Storage (OR) An error log is written to Cloud Logging (OR) A message arrives to Cloud Pub/Sub (OR) A http/https invocation is received
- **Enter Cloud Functions**
  - **Run code in response to events**
    - Write your business logic in Node.js, Python, Go, Java, .NET, and Ruby
    - **Don't worry about servers or scaling or availability (only worry about your code)**
  - **Pay only for what you use**
    - Number of invocations
    - Compute time of the invocations
    - Memory and CPU provisioned
  - **Time Bound** - Default 1 min and MAX 60 minutes (3600 seconds)
  - **2 product versions**
    - Cloud Functions (1st gen): First version
    - Cloud Functions (2nd gen): New version built on top of Cloud Run and Eventarc

# Cloud Functions - Concepts

- **Event** : Upload object to cloud storage
- **Trigger: Respond to event with a Function call**
  - **Trigger** - Which function to trigger when an event happens?
  - **Functions** - Take event data and perform action?
- **Events are triggered from**
  - Cloud Storage
  - Cloud Pub/Sub
  - HTTP POST/GET/DELETE/PUT/OPTIONS
  - Firebase
  - Cloud Firestore
  - Stack driver logging



# Example Cloud Function - HTTP - Node.js

```
const escapeHtml = require('escape-html');

exports.helloHttp = (req, res) => {
  res.send(`Hello ${escapeHtml(req.query.name || req.body.name || 'World')}!`);
};
```


# Example Cloud Function - Pub/Sub - Node.js

```
/**
 * Background Cloud Function to be triggered by Pub/Sub.
 * This function is exported by index.js, and executed when
 * the trigger topic receives a message.
 *
 * @param {object} message The Pub/Sub message.
 * @param {object} context The event metadata.
 */
exports.helloPubSub = (message, context) => {
  const name = message.data
    ? Buffer.from(message.data, 'base64').toString()
    : 'World';


  console.log(`Hello, ${name}!`);
};
```

# Cloud Functions - Remember

- No Server Management: You don't need to worry about scaling or availability of your function
- Cloud Functions automatically spin up and back down in response to events
  - They scale horizontally!
- Cloud Functions are recommended for responding to events:
  - Cloud Functions are NOT ideal for long running processes
    - Time Bound - Default 1 min and MAX 60 minutes(3600 seconds)




# Cloud Run & Cloud Run for Anthos



- **Cloud Run - "Container to Production in Seconds"**
  - Built on top of an open standard - **Knative**
  - **Fully managed serverless platform for containerized applications**
    - ZERO infrastructure management
    - Pay-per-use (For used CPU, Memory, Requests and Networking)
- Fully integrated **end-to-end developer experience:**
  - **No limitations** in languages, binaries and dependencies
  - Easily portable because of **container** based architecture
  - Cloud Code, Cloud Build, Cloud Monitoring & Cloud Logging Integrations
- **Anthos - Run Kubernetes clusters anywhere**
  - Cloud, Multi Cloud and On-Premise
- **Cloud Run for Anthos: Deploy your workloads to Anthos clusters running on-premises or on Google Cloud**
  - Leverage your existing Kubernetes investment to quickly run serverless workloads


# Cloud Functions - Second Generation - What's New?



- **2 Product Versions:**
  - Cloud Functions (1st gen): First version
  - Cloud Functions (2nd gen): New version built on top of Cloud Run and Eventarc
- **Recommended: Use Cloud Functions (2nd gen)**
- **Key Enhancements in 2nd gen:**
  - **Longer Request timeout:** Up to 60 minutes for HTTP-triggered functions
  - **Larger instance sizes:** Up to 16GiB RAM with 4 vCPU (v1: Up to 8GB RAM with 2 vCPU)
  - **Concurrency:** Upto 1000 concurrent requests per function instance (v1: 1 concurrent request per function instance)
  - **Multiple Function Revisions and Traffic splitting** supported (v1: NOT supported)
  - **Support for 90+ event types** - enabled by Eventarc (v1: Only 7)
- DEMO!



# Cloud Functions - Scaling and Concurrency



- **Typical serverless functions architecture:**
  - **Autoscaling** - As new invocations come in, new function instances are created
  - **One function instance handles ONLY ONE request AT A TIME**
  - 3 concurrent function invocations => 3 function instances
    - If a 4th function invocation occurs while existing invocations are in progress, a new function instance will be created
    - HOWEVER, a function instance that completed execution may be reused for future requests
  - (Typical Problem) **Cold Start:**
    - New function instance initialization from scratch can take time
    - (Solution) Configure Min number of instances (increases cost)
- **1st Gen** uses the typical serverless functions architecture
- **2nd Gen** adds a very important new feature:
  - One function instance can handle multiple requests AT THE SAME TIME
    - **Concurrency:** How many concurrent invocations can one function instance handle? (Max 1000)
    - (IMPORTANT) Your function code should be safe to execute concurrently

# Cloud Functions - Deployment using gcloud

- **gcloud functions deploy [NAME]**

- **--docker-registry** (registry to store the function's Docker images)
  - Default - container-registry
  - Alternative - artifact-registry
- **--docker-repository** (repository to store the function's Docker images)
  - Example: (projects/\${PROJECT}/locations/\${LOCATION}/repositories/\${REPOSITORY})
- **--gen2** (Use 2nd gen. If this option is not present, 1st gen will be used)
- **--runtime** (nodejs, python, java,...)
  - Reference - <https://cloud.google.com/functions/docs/runtime-support>
- **--service-account** (Service account to use)
  - 1 GEN - default - App Engine default service account - PROJECT\_ID@appspot.gserviceaccount.com
  - 2 GEN - Default compute service account - PROJECT\_N0-compute@developer.gserviceaccount.com
- **--timeout** (function execution timeout)
- **--max-instances** (function execution exceeding max-instances times out)
- **--min-instances** (avoid cold starts at higher cost)

# Cloud Functions - Deployment using gcloud - 2

```
//Deploy Pubsub Triggered gen2 function from Cloud Storage Bucket
gcloud functions deploy my-pubsub-function \
  --gen2 \
  --region=europe-west1 \
  --runtime=nodejs16 \
  --source=gs://my-source-bucket/source.zip \
  --trigger-topic=my-pubsub-topic
```

- **gcloud functions deploy [NAME]**


- **--source**

- Zip file from Google Cloud Storage (gs://my-source-bucket/my\_function\_source.zip) (OR)
    - Source Repo (https://URL/projects/\${PROJECT}/repos/\${REP0}) (OR)
    - Local file system

- **--trigger-bucket (OR) --trigger-http (OR) --trigger-topic (OR) --trigger-event-filters (ONLY in gen2 - Eventarc matching criteria for the trigger)**

- **--serve-all-traffic-latest-revision (ONLY in gen2)**

# Cloud Functions - Best Practices



- To avoid cold starts, set **min no of instances** (increases cost)
  - Minimize dependencies (loading dependencies increases initialization time)
- Configure **max no of instances** (protect from abnormally high request levels)
- Use Cloud Endpoints (or Apigee or API gateway) for versioning
- Use Cloud Run (& Cloud Functions gen 2) **revisions** for safer releases:
  - Configure which revisions should receive traffic and how much
  - You can rollback to a previous revision, if needed
- Use Secret Manager to securely store secrets (ex: API keys)
- Use **Individual Service Accounts** for each function
  - Grant `roles/cloudfunctions.invoker` role to invoke a cloud function
- Manage dependencies using your language specific tool (npm, pip,...)