



NoSQL Databases

Cloud Datastore and Firestore




- **Datastore** - Highly scalable NoSQL Document Database
 - Automatically scales and partitions data as it grows
 - Recommended for upto a few TBs of data
 - For bigger volumes, BigTable is recommended
 - Supports Transactions, Indexes and SQL like queries (GQL)
 - Does NOT support Joins or Aggregate (sum or count) operations
 - For use cases needing flexible schema with transactions
 - Examples: User Profile and Product Catalogs
 - Structure: Kind > Entity (Use namespaces to group entities)
 - You can export data ONLY from gcloud (NOT from cloud console)
 - Export contains a metadata file and a folder with the data
- **Firestore** = Datastore++ : Optimized for multi device access
 - Offline mode and data synchronization across multiple devices - mobile, IOT etc
 - Provides client side libraries - Web, iOS, Android and more
 - Offers Datastore and Native modes

Understanding Cloud Datastore Best Practices



- Cloud Datastore is a **document store with flexible schema**
 - Recommended for storing things like user profiles
 - Another Use Case: Index for objects stored in Cloud Storage
 - You want to allow users to upload their profile pictures:
 - Store objects (pictures) in Cloud Storage
 - Enable quick search by storing metadata (like ids and cloud storage bucket, object details) in Cloud Datastore
- **Design your keys and indexes carefully:**
 - **Avoid monotonically increasing values as keys**
 - NOT RECOMMENDED - 1, 2, 3, ..., OR "Customer1", "Customer2", "Customer3", ... or timestamps
 - RECOMMENDED - Use `allocateIds()` for well-distributed numeric IDs
 - **Create indexes only if they would be used in queries**
 - For ad hoc queries on large datasets without pre-defined indexes, BigQuery is recommended!
- **Prefer batch operations** (to single read, write or delete operations):
 - More efficient as multiple operations are performed with same overhead as one operation

Cloud BigTable




- **Petabyte scale, wide column NoSQL DB** (HBase API compatible)
 - **Designed for huge volumes** of analytical and operational data
 - IOT Streams, Analytics, Time Series Data etc
 - Handle millions of read/write TPS at **very low latency**
 - **Single row transactions** (multi row transactions NOT supported)
- **NOT serverless**: You need to create a server instance (Use SSD or HDD)
 - Scale horizontally with multiple nodes (No downtime for cluster resizing)
- **CANNOT export data using cloud console or gcloud**:
 - **Either use a Java application** (java -jar JAR export\import) OR
 - **Use HBase commands**
- **Use cbt command line tool to work with BigTable** (NOT gcloud)
 - Ex: cbt createtable my-table

Cloud BigTable - Wide Column Database

Rowid	Column Family 1			Column Family 2			Column Family 3		
	col1	col2	col3	col1	col2	col3	col1	col2	col3
1									
2									
3									

- At the most basic level, each table is a sorted key/value map
 - Each value in a row is indexed using a key - **row key**
 - Related columns are grouped into column families
 - Each column is identified by using column-family:column-qualifier(or name)
- This structure supports high read and write throughput at low latency
 - **Advantages : Scalable to petabytes of data with millisecond responses upto millions of TPS**
- **Use cases** : IOT streams, graph data and real time analytics (time-series data, financial data - transaction histories, stock prices etc)
- **Cloud Dataflow** : Used to export data from BigTable to CloudStorage


Designing BigTable Tables



- **Two things** you should know before starting with Bigtable:
 - What data do you want to store? (format, columns etc)
 - What would your frequently used queries look like (ranked by usage)?
- Design your table: **Cloud Bigtable is a key/value store**
 - **Each table has ONLY ONE index, the row key**
 - **Design your row key** based on your frequently used queries
 - You can have multiple row key segments - Separated by a delimiter (ex: ranga#123456#abcd)
 - Avoid sequential row keys (timestamps or sequential numbers)
 - Include timestamp as part of your row key IF you plan to retrieve data based on the timestamp
 - Use reversed timestamp (Ex: Long.MAX_VALUE - timestamp) if you frequently query recent data
 - Records will be ordered from most recent to least recent
 - **After your design your table:**
 - Test (heavy load for several minutes + one hour simulation) with atleast 30 GB of test data
 - Analyze usage patterns with **Key Visualizer tool** for Cloud Bigtable

Understanding Cloud BigTable Best Practices

- Recommended for streaming IOT & time series data
- Automatically shards data into multiple tablets across nodes in cluster:
 - **Goal 1:** Have same amount of data on each node
 - **Goal 2:** Distribute reads and writes equally across all nodes
 - (REMEMBER) Pre-test with heavy load for a few minutes before you run your tests
 - Gives Bigtable a chance to balance data across your nodes
- Cloud Bigtable supports SSD or HDD storage:
 - **SSD** - For most usecases
 - **HDD** - For large non latency-sensitive data sets of size >10 TB with very very few reads



Understanding Cloud BigTable Best Practices - Replication

- You can create a Cloud Bigtable instance with **more than one cluster to enable replication (Cross Region or Cross Zone)** :
 - Independent copy of data is stored in each cluster (in the zone of the cluster)
 - Bigtable automatically replicates changes
 - Replication improves durability and availability of your data
 - Stores separate copies in multiple zones or regions
 - Can automatically failover between clusters if needed
 - Replication helps you to put data closer to your customers
 - Configure an application profile, or app profile with routing policy of multi-cluster routing
 - Automatically route to nearest cluster in an instance

