


Evolution


Agile > DevOps > SRE

Software development life cycle (SDLC) - Waterfall



- **Software development in multiple long phases:**
 - Requirements
 - Design
 - Implementation
 - Testing
 - Deployment

Software development life cycle (SDLC) - Spiral



- **Software development in smaller iterations:**

- Start
- Iteration 1
- Iteration 2
- ...
- ...


Software development life cycle (SDLC) - Agile

- Principles

- Individuals and interactions over processes and tools
- Working software over comprehensive documentation
- Customer collaboration over contract negotiation
- Responding to change over following a plan
- Now there are 12 principles
(<https://agilemanifesto.org/principles.html>)

- Agile is recommended for most software development:

- BUT add a bit of rigidity from waterfall model for critical safety software (Flight navigation software, Medical devices software etc)

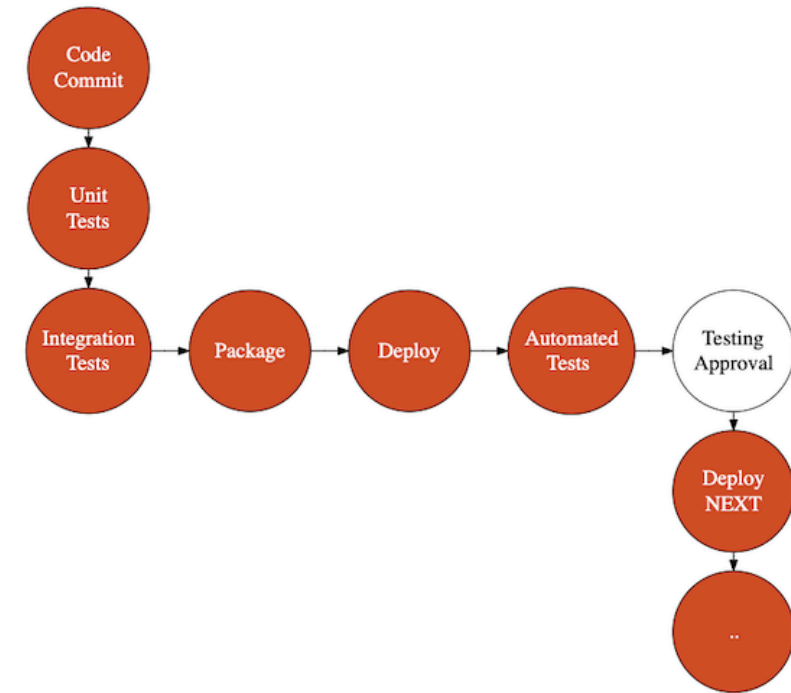




- Getting Better at **"Three Elements of Great Software Teams"**
 - **Communication** - Get teams together
 - **Feedback** - Earlier you find a problem, easier it is to fix
 - **Automation** - Automate testing, infrastructure provisioning, deployment, and monitoring

DevOps - CI, CD

- **Continuous Integration**
 - Continuously run your tests and packaging
- **Continuous Deployment**
 - Continuously deploy to test environments
- **Continuous Delivery**
 - Continuously deploy to production



DevOps - CI CD - Recommended Things to Do

- **Static Code Analysis**

- Lint, Sonar
- Including Static Security Checks (Source Code Security Analyzer software like Veracode or Static Code Analyzer)

- **Runtime Checks**

- Run Vulnerability Scanners (automated tools that scan web applications for security vulnerabilities)

- **Tests**

- Unit Tests (JUnit, pytest, Jasmine etc)
- Integration Tests (Selenium, Robot Framework, Cucumber etc)
- System Tests (Selenium, Robot Framework, Cucumber etc)
- Sanity and Regression Tests

DevOps - CI, CD Tools

- **Cloud Source Repositories**: Fully-featured, private Git repository
 - Similar to Github
- **Container Registry**: Store your Docker images
- **Jenkins**: Continuous Integration
- **Cloud Build**: Build deployable artifacts (jars or docker images) from your source code and configuration
- **Spinnaker**: Multi-cloud continuous delivery platform
 - Release software changes with high velocity and confidence
 - Supports deployments to Google Compute Engine, Google Kubernetes Engine, Google App Engine and other cloud platforms
 - Supports Multiple Deployment Strategies



Container
Registry




Cloud Source
Repositories

DevOps - Infrastructure as Code



- Treat infrastructure the same way as application code
- Track your infrastructure changes over time (version control)
- Bring repeatability into your infrastructure
- Two Key Parts
 - Infrastructure Provisioning
 - Provisioning compute, database, storage and networking
 - Open source cloud neutral - Terraform
 - GCP Service - Google Cloud Deployment Manager
 - Configuration Management
 - Install right software and tools on the provisioned resources
 - Open Source Tools - Chef, Puppet, Ansible and SaltStack


Google Cloud Deployment Manager - Introduction



- Lets consider an example:
 - I would want to create a new VPC and a subnet
 - I want to provision a Load balancer, Instance groups with 5 Compute Engine instances and an Cloud SQL database in the subnet
 - I would want to setup the right Firewall
- AND I would want to create 4 environments
 - Dev, QA, Stage and Production!
- Deployment Manager can help you do all these with a simple (actually NOT so simple) script!

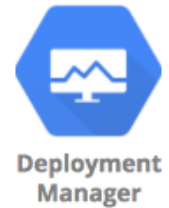
Google Cloud Deployment Manager - Advantages

- Automate deployment and modification of Google Cloud resources in a controlled, predictable way
 - Deploy in multiple environments easily!
- Avoid configuration drift
- Avoid mistakes with manual configuration
- Think of it as version control for your environments
- **Important Note** - Always modify the resources created by Deployment Manager using Deployment Manager



Google Cloud Deployment Manager

- All configuration is defined in a simple text file - **YAML**
 - I want a VPC, a subnet, a database and ...
- Deployment Manager understands dependencies
 - Creates VPCs first, then subnets and then the database
- **(Default) Automatic rollbacks on errors (Easier to retry)**
 - If creation of database fails, it would automatic delete the subnet and VPC
- Version control your configuration file and make changes to it over time
- **Free to use - Pay only for the resources provisioned**
 - Get an automated estimate for your configuration



Cloud Deployment Manager - Example

```
- type: compute.v1.instance
  name: my-first-vm
  properties:
    zone: us-central1-a
    machineType: <<MACHINE_TYPE>>
    disks:
      - deviceName: boot
        type: PERSISTENT
        boot: true
        autoDelete: true
        initializeParams:
          sourceImage: <<SOURCE_IMAGE>>
    networkInterfaces:
      - network: <<NETWORK>>
        # Give instance a public IP Address
        accessConfigs:
          - name: External NAT
            type: ONE_TO_ONE_NAT
```

Cloud Deployment Manager - Terminology

- **Configuration** file: YAML file with resource definitions for a single deployment
- **Templates**: Reusable resource definitions that can be used in multiple configuration files
 - Can be defined using:
 - Python (preferred) OR
 - JinJa2 (recommended only for very simple scripts)
- **Deployment**: Collection of resources that are deployed and managed together
- **Manifests**: Read-only object containing original deployment configuration (including imported templates)
 - Generated by Deployment Manager
 - Includes fully-expanded resource list
 - Helpful for troubleshooting

Cloud Marketplace (Cloud Launcher)

- Installing custom software might involve setting up multiple resources:
 - Example: Installing WordPress needs set up of compute engine and a relational database
- How do you simplify the set up of custom software solutions like Wordpress or even more complex things like SAP HANA suite on GCP?
- **Cloud Marketplace:** Central repo of easily deployable apps & datasets
 - **Similar to App Store/Play Store for mobile applications**
 - You can search and install a complete stack
 - Commercial solutions - SAP HANA etc
 - Open Source Packages - LAMP, WordPress, Cassandra, Jenkins etc
 - OS Licenses: BYOL, Free, Paid
 - Categories: Datasets/Developer tools/OS etc
 - When selecting a solution, you can see:
 - Components - Software, infrastructure needed etc
 - Approximate price

Site Reliability Engineering (SRE)

- DevOps++ at Google
- SRE teams focus on every aspect of an application
 - availability, latency, performance, efficiency, change management, monitoring, emergency response, and capacity planning
- Key Principles:
 - Manage by Service Level Objectives (SLOs)
 - Minimize Toil
 - Move Fast by Reducing Cost of Failure
 - Share Ownership with Developers



Google Cloud

Site Reliability Engineering (SRE) - Key Metrics

- **Service Level Indicator(SLI):** Quantitative measure of an aspect of a service
 - Categories: availability, latency, throughput, durability, correctness (error rate)
 - Typically aggregated - "Over 1 minute"
- **Service Level Objective (SLO) -** SLI + target
 - 99.99% Availability, 99.999999999% Durability
 - Response time: 99th percentile - 1 second
 - Choosing an appropriate SLO is complex
- **Service Level Agreement (SLA):** SLO + consequences (contract)
 - What is the consequence of NOT meeting an SLO? (Defined in a contract)
 - Have stricter internal SLOs than external SLAs
- **Error budgets: (100% – SLO)**
 - How well is a team meeting their reliability objectives?
 - Used to manage development velocity

Site Reliability Engineering (SRE) - Best Practices



Google Cloud

- **Handling Excess Loads**

- **Load Shedding**

- API Limits
 - Different SLAs for different customers
 - Streaming Data
 - If you are aggregating time series stream data, in some scenarios, you can drop a part of data

- **Reduced Quality of Service**

- Instead of talking to a recommendations API, return a hardcoded set of products!
 - Not always possible:
 - Example: if you are making a payment

- **Avoiding Cascading Failures**

- **Plan to avoid thrashing**

- Circuit Breaker
 - Reduced Quality of Service

Site Reliability Engineering (SRE) - Best Practices - 2

- **Penetration Testing (Ethical Hacking)**

- **Simulate an attack** with the objective of finding security vulnerabilities
- Should be authorized by project owners
- No need to inform Google
 - Ensure you are only testing your projects and are in compliance with terms of service!
- Can be white box (Hacker is provided with information about infrastructure and/or applications) or black box (No information is provided)

- **Load Testing (JMeter, LoadRunner, Locust, Gatling etc)**

- **Simulate real world traffic as closely as possible**
- **Test for spiky traffic - suddenly increases in traffic**



Google Cloud

Site Reliability Engineering (SRE) - Best Practices - 3

- **Resilience Testing** - "How does an application behaves under stress?"
- **Resilience** - "Ability of system to provide acceptable behavior even when one or more parts of the system fail"
- **Approaches:**
 - **Chaos Testing (Simian Army)** - cause one or more layers to fail
 - "unleashing a wild monkey with a weapon in your data center to randomly shoot down instances and chew through cables"
 - Add huge stress on one of the layers
 - **Include network in your testing** (VPN, Cloud Interconnect etc..)
 - Do we fall back to VPN if direct interconnect fails?
 - What happens when internet is down?
 - **Best Practice: DiRT - disaster recovery testing at Google**
 - Plan and execute outages for a defined period of time
 - Example: Disconnecting complete data center



Google Cloud