


Release Management

Release Management

- **Goals: vary from app to app**
 - Zero Downtime
 - Only one version live at a time
 - Minimize Costs (and infrastructure needed)
 - Test with production traffic before going live
- **Best Practices:**
 - Small incremental changes
 - Automation (as much as possible)
 - Handling problems with new releases:
 - Analyze logs and metrics from Cloud Monitoring and Logging
 - Rollback to previous release and try replicating the problem in other environments



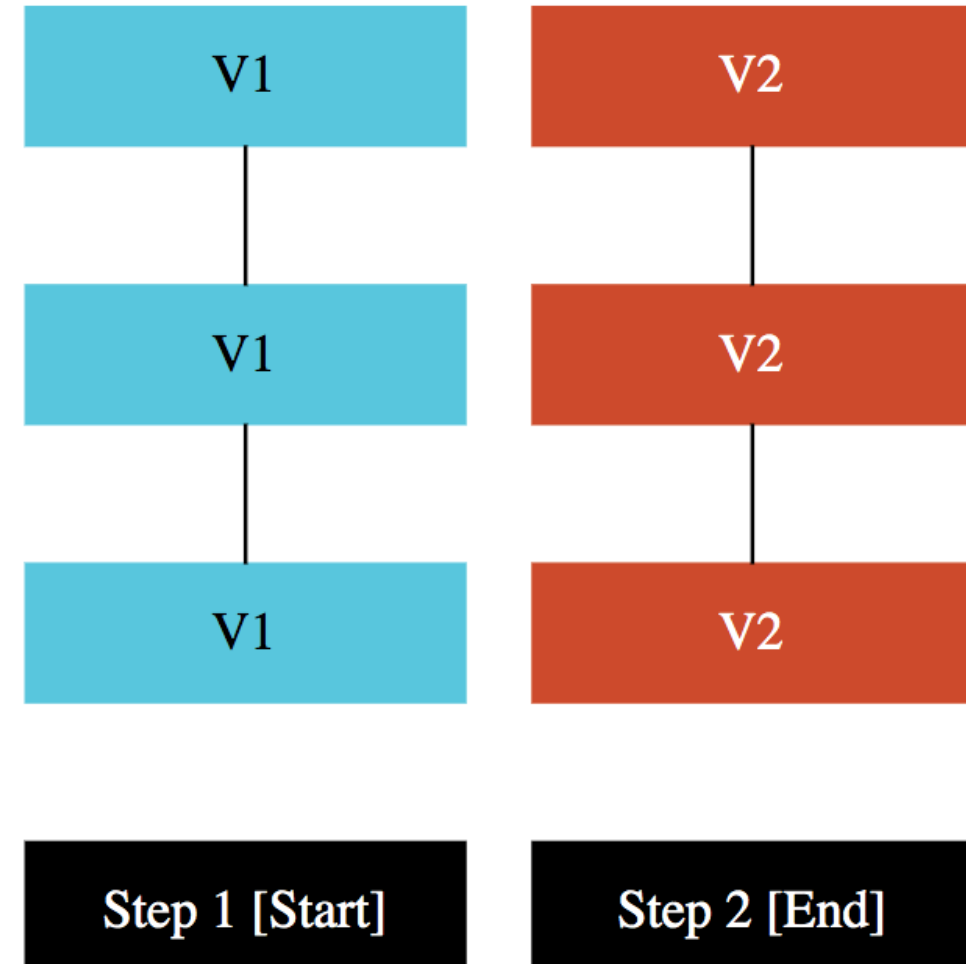
Deployment Approach : Recreate

- **Approach:**

- Terminate Version 1
- Roll out Version 2

- **Characteristics:**

- App down during the release
- Rollback needs redeployment
 - AND more downtime
- Cost effective and Fast
 - BUT disruptive
- Avoid need for backward compatibility (data and applications)




Deployment Approach : Canary

- **Approach:**

- Step 1: V2 rolled out to a subset of instances
- Step 2: On successful testing, V2 rolled out to all instances
 - OR V2 is rolled back in case of failure


- **Characteristics:**

- Fast
- Zero downtime
- No extra infrastructure
- Minimizes impact to users (in case of release failures)
- Needs Backward compatibility (data and applications)



Testing Approach : A/B Testing

- Use case: You want to see if users like a feature!
- Approach:
 - Step 1: V2 (with new feature) rolled out to a subset of users
 - Step 2: On successful testing, V2 rolled out to all users
 - OR we go back to V1 in case users don't like the feature!
- Characteristics:
 - Gives the ability to test if your users like a feature




Deployment Approach : Rolling

- **Approach:**

- **Step 1:** V2 rolled out to a percentage of instances (Example window size: 5%)
- **Step 2..N:** V2 gradually rolled out to rest of the instances (Example: 5% at a time)

- **Characteristics:**

- Slow
- Zero downtime
- Needs automation and additional setup
- No extra infrastructure
- Minimizes impact to users (in case of release failures)
- Needs Backward compatibility (data and applications)



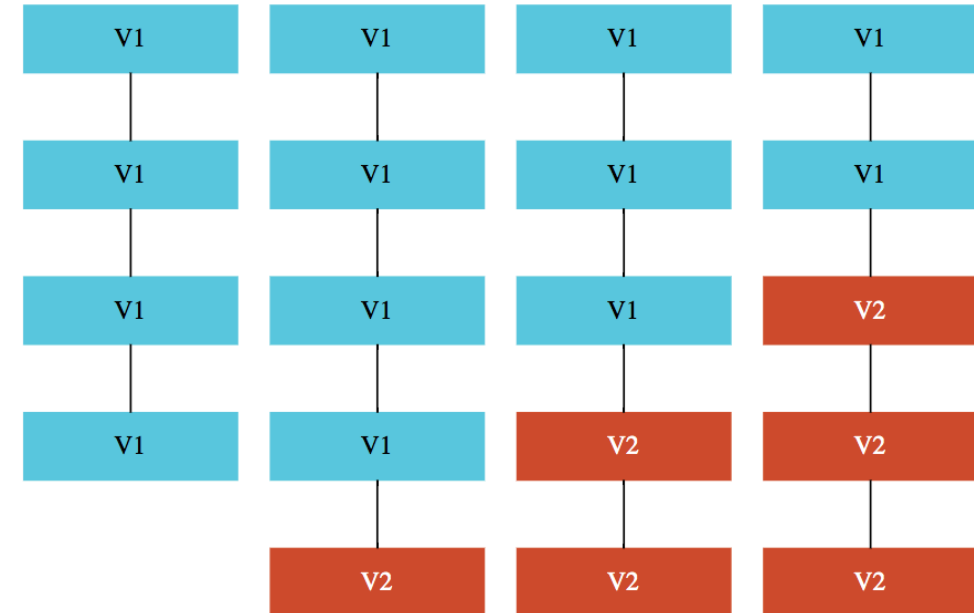
Deployment Approach : Rolling with Additional Batch

- **Approach:**

- **Step 1:** Additional batch of new instances are created with V2 (Example: 5%)
- **Step 2..N:** V2 gradually rolled out to the instances batch by batch (Example: 5% at a time)

- **Characteristics:**

- Same as Rolling Deployment except for:
 - Needs Little bit of extra infrastructure
 - ZERO reduction in number of instances handling user requests




Deployment Approach : Blue Green

- Approach:

- Step 1: V1 is Live
- Step 2: Create (or replicate) a parallel environment with V2
- Step 3: Switch all traffic from V1 to V2 (and remove V1 Environment)

- Characteristics:

- Instant
- Zero Downtime
- Easy Rollback
- Needs additional infra (during the release)
- ZERO reduction in available capacity
- Needs Backward compatibility (data and apps)




Testing Approach : Shadow

- Approach:

- Step 1: V1 is Live
- Step 2: Create (or replicate) a parallel environment with V2
 - Mirror traffic to V1 and V2
- Step 3: Switch all traffic from V1 to V2 (and remove V1 Environment)

- Characteristics:

- Zero production impact: Test V2 with real production traffic before releasing
 - You can also capture and replay live production traffic
- Complicated : You don't want double payments (might need stubbing)
- Needs a lot of additional infrastructure



Deployment Approaches - Managed instance group (MIG)

Option	Details
Rolling Release	<i>gcloud compute instance-groups managed rolling-action start-update my-mig --version=template=v2-template</i> --max-surge=5 or 10% (Max instances updated at a time) --max-unavailable=5 or 10% (Max instances that can be down during update)
Canary Release	<i>gcloud compute instance-groups managed rolling-action start-update my-mig --version=template=v1-template</i> --canary-version=template=v2-template,target-size=10%
Blue Green Deployment	Create a new MIG and make manual adjustments to Load Balancer backends as needed

App Engine - Releasing New Versions

Option	Details
Deploy & shift all traffic at once	<code>gcloud app deploy</code>
Deploy v2 without shifting traffic	<code>--no-promote</code>
Shift traffic to V2 all at once	<code>gcloud app services set-traffic s1 --splits V2=1</code>
Gradual Migration	Add <code>--migrate</code> option to previous command
A/B testing	<code>gcloud app services set-traffic s1 --splits=v2=.5,v1=.5</code>

GKE - Releasing New Versions

Option	Details
Recreate	Set <code>strategy > type</code> on Deployment to <code>Recreate</code> Use 'kubectl set image deployment' OR Update deployment YAML to perform deployment
Rolling Update	Set <code>strategy > type</code> on Deployment to <code>RollingUpdate</code> Use 'kubectl set image deployment' or Update deployment YAML to perform deployment Configure <code>maxSurge</code> and <code>maxUnavailable</code> to control deployment
Blue Green Deployment	Create New Deployment. Control traffic using Ingress (or Service)
Canary Deployment	Needs Service Mesh like Istio